

**INVESTIGATING THE ROLE OF PREDICTIVE
REPRESENTATIONS IN IMPLICIT EVENT
BOUNDARIES, STATISTICAL LEARNING, AND
CATEGORIZATION.**

A Dissertation Presented

by

TEJAS SAVALIA

Submitted to the Graduate School of the
University of Massachusetts Amherst in partial fulfillment
of the requirements for the degree of

DOCTOR OF PHILOSOPHY

September 2024

Psychological and Brain Sciences

© Copyright by Tejas Savalia 2024

All Rights Reserved

**INVESTIGATING THE ROLE OF PREDICTIVE
REPRESENTATIONS IN IMPLICIT EVENT
BOUNDARIES, STATISTICAL LEARNING, AND
CATEGORIZATION.**

A Dissertation Presented

by

TEJAS SAVALIA

Approved as to style and content by:

Andrew Cohen, Chair

Jeffrey Starns, Member

Youngbin Kwak, Member

Meghan Huber, Member

Maureen Perry-Jenkins, Chair
Psychological and Brain Sciences

DEDICATION

To those little lost sheep.

ACKNOWLEDGMENTS

Thanks to all those fine shepherds. Not to mention all the great border collies and suchlike fine animals.

ABSTRACT

INVESTIGATING THE ROLE OF PREDICTIVE REPRESENTATIONS IN IMPLICIT EVENT BOUNDARIES, STATISTICAL LEARNING, AND CATEGORIZATION.

SEPTEMBER 2024

TEJAS SAVALIA

B.E., GUJARAT TECHNOLOGICAL UNIVERSITY

M.S., INTERNATIONAL INSTITUTE OF INFORMATION TECHNOLOGY,
HYDERABAD

M.S., UNIVERSITY OF MASSACHUSETTS AMHERST

Ph.D., UNIVERSITY OF MASSACHUSETTS AMHERST

Directed by: Professor Andrew Cohen

Sheep like grass. Why? Let me tell you. Sheep are ruminants, like cattle, deer, and horses. They have stomachs that are specialized...

TABLE OF CONTENTS

	Page
ACKNOWLEDGMENTS	v
ABSTRACT	vi
LIST OF TABLES	ix
LIST OF FIGURES	x
CHAPTER	
1. INTRODUCTION	1
1.1 Scope of this dissertation	4
1.2 Format of this dissertation	4
2. QUALITY OF ENVIRONMENTAL EXPOSURE MODULATES STATISTICAL LEARNING	6
2.1 Introduction	6
2.1.1 Representations of Temporal Context	8
2.1.1.1 Successor Representation	8
2.1.2 Temporal Context Model	9
3. CHAPTER 4: CATEGORY LEARNING THROUGH TEMPORAL ABSTRACTION	10
3.1 (Partly polished) Introduction	10
3.1.1 Participants	14
3.1.2 Materials	14
4. DISCUSSION	15

BIBLIOGRAPHY	16
--------------------	----

LIST OF TABLES

Table

Page

LIST OF FIGURES

Figure

Page

CHAPTER 1

INTRODUCTION

We experience a constant stream of sensory information from the moment we are born. Our brain parses this information and slowly learns to extract meaning from it. From recognizing the mother's scent as a survival instinct to formulating complex plans to defeat a board game opponent, our brain extracts meaning from our surroundings, considers prior experience and the current state of the world, and makes decisions to interact accordingly. My key question in this dissertation is: How do we learn to extract meaning from our surroundings?

Extracting meaningful information from our experience is beneficial to our functioning and survival. When approached by a cheetah in a forest, we do not wait to evaluate the exact number of spots on its skin before deciding to run. Such abstractions and formation of heuristics allow us to process naturally complex surroundings in quick time and act accordingly.

When extracting meaning from our surroundings, we often (need to) ignore minute details and abstract out towards a coherent thought of our surroundings. The extreme example above aside, we observe such abstractions in almost all day-to-day activities. Imagine someone asking you about the events during the day before reading this manuscript. Perhaps you are reading it on your office computer and you recall a sequence of events starting from waking up, to getting ready, driving or taking public transit to work, getting your morning coffee and breakfast in with an email check before turning your attention to this manuscript. Each event described above combines several sub-events that are abstracted away in such a verbal recall and de-

scription. For example, getting ready involves several steps from brushing your teeth, showering, and wearing your work outfits. Each sub-event can be further thought of as an abstraction from sub-sub-events – brushing your teeth is a combination of putting toothpaste on the brush head, the physical act of brushing, followed by rinsing. While we perform each act continuously in time, our recall (and by extension, representation in memory) of these past events is discrete, segmented, and abstracted. The meaningfully separatable chunks of a continuous stream of events help in storage in (and retrieval from) memory.

There is often agreement on what it means for a chunk and for boundaries defining such chunks to be ‘meaningful’. In the example above, it is reasonable to argue that brushing teeth, showering, and putting on clothes are three distinct activities. Furthermore, even when the true transitions between these activities are continuous and seamless to a independent, naive observer, the boundaries between these events are perceptually meaningful. What aspect of the environment dictates this agreement about the points at which we segment events and what is special about the properties of the events between those points that lead to distinct representations in memory?

One could argue that these events that occur at different points in time also occur at different spatial locations, thus providing different contexts and hence separate representations in our brain to be considered distinct. Transitioning from one (temporal) event to another can thus be akin to transitioning from one room of the house to another. However, it is almost impossible to decouple temporal events from spatial events assuming a causality from spatial segmentation to temporal segmentation for any spatially experienced distinct event is also a temporally experienced distinct event. Instead, arguing that temporally distinct events lead to a spatially distinct representation can provide a more encompassing explanation of distinct representations for events in distinct spatial *and* temporal contexts. One could similarly argue that the formation of “meaningful” chunks is through perceptual differences between

the events we experience. While lower-level perceptual experiences are indeed often different for different events, the mapping of perceptual differences onto segmented events is arbitrary and not a *sufficient* condition. For example, eating an apple is recalled as eating an apple regardless of whether it has a green leaf added to its top. Perceptual distinctions are not enough to determine whether events are represented distinctly. What then is the key mechanism that leads to events being segmented?

In this dissertation, I argue that the primary reason and mechanism through which we segment events is based on temporal contingencies of various sub-events that encompass an event. Specifically, we recall being in the kitchen as different from being in the bedroom because we have a coherent set of experiences in the kitchen that are distinct from a coherent set of experiences in the bedroom. For an infant forming knowledge of the world, a kitchen while perceptually distinct from a bedroom, is not meaningfully different. With experience and observation of the functions within these spatially (and perceptually) distinct locations, the child slowly develops distinct representations of the two rooms.

This dissertation focuses on understanding temporal contingencies' role in event segmentation, and by extension, general pattern extraction. I argue that even *without* any spatial or perceptual information that may aid us in separating events in memory, we can use temporal coherence to experience separate events and abstract information to aid higher-level cognition. Specifically, I investigate the parsing of a continuous stream of information into discrete chunks in three ways:

- The possible algorithmic processes that naturally lead to such segmentation and the impact of environmental properties in aiding this abstraction.
- The properties of the temporal boundaries when such temporal abstraction occurs naturally and implicitly.

- The role of temporal events separated by implicit temporal boundaries in forming abstractions.

1.1 Scope of this dissertation

The human brain is a complex machine – millions of neurons act as computational units and combine in specific ways to form a functioning human being. These neurons come together to implement several levels of function from lower-level automatic perception to higher-level planning and conscious thought. This dissertation does **not** focus on these implementational-level mechanisms of cognition. Rather, it focuses on *algorithmic* computations that neurons may, collectively, implement that lead to us acquiring patterns in the environment around us.

Analyses in this dissertation use several models of cognition in investigating the role of implicit statistical learning. In most cases, the focus of this dissertation is **not** to evaluate the validity of these models. Indeed, most models of cognition are wrong; but are useful and I use several such models to evaluate specific aspects of how we acquire patterns.

1.2 Format of this dissertation

In the rest of this dissertation, I present three lines of studies investigating the role of implicit temporal boundaries in cognition. In Chapter 2, I present an algorithmic framework that naturally leads to a representation of separable events without the need to rely on explicit properties of the experienced events. I also show that this framework allows for a distinct representation of event boundaries as special events. In Chapter 3, I present work comparing the properties of these event boundaries which are operationalized implicitly (i.e. through no perceptually special information) with event boundaries as they have been studied in prior literature which are operationalized explicitly. Finally, in Chapter 4, I present work investigating the role of these

implicitly operationalized event boundaries in categorization to serve as a gateway for understanding higher-order cognition in the context of temporal segmentation and pattern acquisition.earned.

CHAPTER 2

QUALITY OF ENVIRONMENTAL EXPOSURE MODULATES STATISTICAL LEARNING

2.1 Introduction

Imagine you just moved to the United States and are visiting Target for the first time. Perhaps since you just moved in, your first goal is to furnish your apartment. You look around at the entrance, and navigate your way to the furniture section perhaps while taking a few false turns on the way, buy the stuff you need, pay and leave. The next time you visit for, for example, groceries and produce. You visit again for sporting goods and then again for gifts for your friends. An year later, all such subset of visits through the Target store makes you an expert in knowing the specific route to the section you need to visit. What algorithmic mechanisms allow us to build such expertise to create connections in an explored environment even when during no single visit, you explored all possible connections between regions of the store?

We are able to build up a map of the environment without being exposed to the full extent of it in a single iteration simply based on local exposures to it. Such building is often implicit – you just know that in order to go to the furniture area, you need to pass through the gifts even when you did not explicitly explore this connection before. In this chapter I explore some algorithmic mechanisms that may lead to our ability to acquire large-scale structural knowledge of the environment based on local exposures.

The effect of local exposure in acquiring structural knowledge of the environment have been explored in several areas of cognitive psychology through artificial grammar

and language learning [21, 40, 1, 11], visual statistical learning [12, 44, 4], or motor sequence learning [2, 34, 7, 17]. In recent work, (implicit) acquisition of higher order knowledge of the environment from lower order exposure is studied through structured graph based transitions between stimuli [41, 20, 17, 27, 29, 18]. For example [41] show that when asked to respond to arbitrary stimuli arranged in a temporally graph-modular structure, participants often parse the edges that connect two modules as ‘natural breaks’ even when their local exposure does not distinguish between the cross-module and within-module edges.

More commonly, global-scale structure acquisition has been attributed to response time measurements. Earlier work in serial reaction time tasks [34, 7] shows that breaking an implicitly learned motor sequence leads to slower reaction times. The slowed reaction times when crossing the between-module edges have also been shown in recent work on statistical learning in modular graph structures [17, 27, 20, 18, 19, 29].

This slowdown across module edges appears to be mediated by the nature of the walk experienced across the community structure where random and Eulerian walk (a walk where each edge of the graph is visited exactly once before repeats) experiences continue to show this slowdown whereas a Hamiltonian walk (a walk where each node of the graph is visited exactly once before repeats) experience does not [20]. Thus it appears that the kind of experience through the graph alters the knowledge of underlying statistical patterns. Similarly, the topographical structure of a graph in motor skill learning tasks also appears to alter structural knowledge [28, 29, 27] where modular graphs like in Figure ?? produce the largest dip in reaction times when responding to boundary items.

Why do we slow down at boundary nodes that lead to the adjacent module even when the local probability of that particular transition is the same as any other transitions? Understanding this particular property of human behavior may provide

deeper insights into what leads to global-scale structure acquisition; after all the only difference between the boundary node and other non-boundary nodes is in context of the global structure of the graph. Event boundary literature (where boundaries are typically operationalized through explicit changes in context) suggests that boundaries alter the predictability of future events and this predictability leads to event segmentation [50, 8]. Thus, in implicitly operationalized boundaries such as in serial reaction time tasks, the slowdown at the boundary node may imply a similarly increased uncertainty at boundary nodes leading to slowed responses. Prior work aimed at understanding human representation of graph structures indeed points to an the ‘cross-entropy’ between a learner’s estimate of the transition probability and the true transition probability of the environment [28, 27, 29].

In particular, [29] show that algorithms of contextual representations such as the Successor Representation (SR) model in Reinforcement Learning [10, 33, 14] or the associative learning based Temporal Context Model (TCM) can naturally lead to an increased cross-entropy for cross cluster transitions relative to within cluster transitions in modular graphs. In this work, by using the framework of cross-entropy to estimate reaction times in a modular graph we aim to 1) Experimentally test the predictions of these two models when exposure through the modular graph structure of is partial and 2) Provide evidence in favor of one of the two models of representation.

2.1.1 Representations of Temporal Context

2.1.1.1 Successor Representation

The Successor Representation (SR) model of reinforcement learning has been used as a model to understand the generalization of reinforcement learning behavior in large action spaces [10]. In recent work, the SR model has also been shown to be a reliable model for explaining human decision-making behavior in multi-step environments. The model accurately predicted that humans are worse at adapting to

changes in the transition probability of a learned environment than to changes in the end-point rewards [33]. There has been further evidence of SR being represented in the Hippocampal cells which represent space [14, 43].

Briefly, the SR model represents each state in the actionable space as a predictive representation matrix. For an environment of N discrete states, the SR matrix M of size $(N \times N)$ maintains expected future visits to a given state from each state. Specifically, element $M_{i,j}$ of the matrix represents the expected future visits to state j from state i . This transition matrix is learned over time based on the temporal difference error learning rule. For example, consider at a given point in time, t , an agent maintaining the SR matrix is in state i . The agent now moves to state j out of the possible N states. The i^{th} row of the SR matrix is updated as follows:

$$M_{i,j} \leftarrow \hat{M}_{i,j} + \alpha[\delta(s_{t+1}, j) + \gamma * \hat{M}_{s_{t+1},j} - \hat{M}_{s_t,j}] \quad (2.1)$$

where $\delta(.,.)$ equates to 1 if both arguments are equal otherwise it equates to 0. Thus, the matrix increases the probability of visiting a state j from state i if state j is visited in the current experience and it decreases the probability of visiting all other states from state i . Parameter α is a learning rate parameter that determines how much of the previous estimate of visiting state j from i is factored into the current update. Parameter γ is a future discount parameter that dictates how much in the future the agent sees – specifically, a higher value of γ indicates future visitations to state j are weighed high in the current update.

2.1.2 Temporal Context Model

The Temporal Context Model (TCM) was devised to explain the primacy and recency effects in human recall and recognition memory [16]. The TCM model assumes that the items maintain a temporal context as they get encoded thus allowing items presented close to the previous items to share such temporal context.

CHAPTER 3

CHAPTER 4: CATEGORY LEARNING THROUGH TEMPORAL ABSTRACTION

3.1 (Partly polished) Introduction

We naturally categorize items we encounter daily for ease of storage, processing, and decision-making. For example we know instinctively that regardless of shape and form, all lamps form a 'lamp' category based on its function. Several factors determine how we categorize items. Depending on the complexity of rules that determine categories, some categorizations are easier than others [42, 37]. Category variability can modulate how often exemplars are classified into that categories [9].

The order of presentation items in category learning tasks has been shown to be an important factor in how category diagnostic features are learned. In particular, when items are presented in a blocked categorical design, participants seem to learn the similarities between the same category items. On the other hand, when items are presented as an interleaved design, participants seem to focus more on learning the features that differentiate the underlying categories [6]. As a result of order-dependent differing focus on category diagnostic features, interleaved presentations seem to benefit general category learning. In most prior category-learning tasks assessing order of presentation effects, participants are explicitly asked to learn the underlying categories. There appear to be clear differences when participants focus on learning categories based on how exemplars of these categories are presented [22, 23, 49, 47, 5, 6]. In this article, we investigate the effects of order of presentation when category learning is implicit.

One primary focus on category learning through order of presentation is comparing blocked or interleaved exemplar presentations. For example, [22] showed participants paintings made by two different painters. The order of presentation during exposure was modulated to either be blocked (paintings of one artist shown together followed by the second artist) or interleaved (paintings made by both artists were mixed). When presented with new paintings, and asked which of the two studied artists made them, participants who were exposed to the interleaved format were found to be more accurate at guessing the creator. Category learning also improved for interleaved presentation compared to blocked presentation when tested on items where relevant category features were visually occluded [49]. When three-year-old children are tested on the generalization of category-specific features, they appear to benefit from the spaced study of exemplars as compared to a blocked [47]. By modulating the similarity of presented items, interleaved presentation was found to be better than blocked presentation design on generalization performance particularly when learned exemplars were more visual [22, 5].

Interleaved presentation has been theorized to improve in category induction because of context-based variability during encoding [15]. Particularly, for each presented item, an observer will store both the item-specific features along with the context in which the item is encoded. During interleaved presentation, a category diagnostic feature gets encoded under different contexts. Thus, that diagnostic feature will be recalled when tested on novel category items within that context.

Two theories have been proposed to explain this discriminability-based advantage of interleaving. According to the attention attenuation account, when categories are blocked, participants may think that they have learned the relevant category features after viewing a few items and stop paying attention to additional exemplars of the [23]. On the other hand, according to the discrimination account, the interleaved presentation allows participants to directly compare the differences between exemplars

of different categories that are presented close to each other thereby highlighting these differences [22]. In a direct test [48] found that when participants were shown pairs of exemplars, each belonging to a different category, the interleaving benefit was magnified compared to when they were presented as single items. The authors posit that showing pairs of exemplars would enable participants to carefully study and infer distinctions between category features and hence improve categorization performance. Furthermore, the authors find evidence against the attention attenuation theory by observing that classification performance did not differ as a function of the position in which the exemplar was presented in a stream.

This benefit of interleaved presentation is shown to be modulated by the ‘level’ at which categorization occurs. For example, when [30] modulated exposure time to individual exemplars along with order of exposure, they found that interleaved presentation was no longer beneficial under short exposure conditions particularly when participants were asked to make a more abstract, ‘super-ordinate level categorization. On the other hand, when exemplars were presented in a blocked format, a lower, ‘basic’ level categorization was hindered. Thus, category knowledge through order of presentation can be modulated by the level of categorization participants are asked to produce.

It is clear that order of presentation of categories matters during explicit category learning. **However, the effect of such order of presentation has not been investigated when category learning is implicit.** Indeed recent work shows that participants do acquire category knowledge that when presented implicitly instead of being explicitly asked to learn categories. [45] found that assessed on category knowledge, participants appeared to learn category structures without being explicitly instructed to do so. This category knowledge was modulated by the strength of association of the category diagnostic features. [46] later found that when presented

with implicit feature-based categories during a cover task, participants were sensitive towards category diagnostic knowledge.

More evidence for incidental category learning comes from auditory cognition. [13] found that participants were sensitive to audio categories learned implicitly as measured by increased reaction times when audio-category-to-response mapping was altered. Incidental category knowledge is further modulated by the sampling category distributions from which exemplars are drawn. [38] show that probabilistic sampling of exemplars leads to weaker category learning compared to deterministic sampling. Incidental category learning can be further enhanced by task-relevant and disrupted by task-irrelevant feature-to-category mappings [39]. Incidental learning may also be disadvantaged compared to supervised intentional learning when categories are non-linearly separable [25].

Thus, while implicit category learning appears to be consistent and dependent on several aspects of the underlying categories, unlike explicit category learning, it is unclear whether implicit category learning enjoys the same advantage when category exemplars are presented in an interleaved vs. a blocked design. Most implicit categorization tasks involve manipulation of features as opposed to manipulation of the temporal order of exposure. In this article, we explore the effects of temporal co-occurrences to manipulate

3.1.0.0.1 Other stuff to incorporate: Unsupervised category learning [3], participants could infer rules based on correlating features without being explicitly asked to categorize during exposure.

Category-related items were recognized better when presented close to each other than when category-unrelated items were [31]. In [32], people unsupervised sorted stimuli by a single dimension, ignoring the family resemblance structure.

SUSTAIN [26] seems to explain all these unsupervised category learning phenomena. ALCOVE [24] provides for error based diagnostic feature attention learning in GCM [36, 35].

3.1.1 Participants

3.1.2 Materials

CHAPTER 4

DISCUSSION

BIBLIOGRAPHY

- [1] Aslin, Richard N, and Newport, Elissa L. Statistical learning: From acquiring specific items to forming general rules. *Current directions in psychological science* 21, 3 (2012), 170–176.
- [2] Baldwin, Dare, Andersson, Annika, Saffran, Jenny, and Meyer, Meredith. Segmenting dynamic human action via statistical structure. *Cognition* 106, 3 (2008), 1382–1407.
- [3] Billman, Dorrit, and Knutson, James. Unsupervised concept learning and value systematicity: A complex whole aids learning the parts. *Journal of Experimental Psychology: Learning, Memory, and Cognition* 22, 2 (1996), 458.
- [4] Brady, Timothy F, and Oliva, Aude. Statistical learning using real-world scenes: Extracting categorical regularities without conscious intent. *Psychological science* 19, 7 (2008), 678–685.
- [5] Carvalho, Paulo F, and Goldstone, Robert L. Putting category learning in order: Category structure and temporal arrangement affect the benefit of interleaved over blocked study. *Memory & cognition* 42 (2014), 481–495.
- [6] Carvalho, Paulo F, and Goldstone, Robert L. The sequence of study changes what information is attended to, encoded, and remembered during category learning. *Journal of Experimental Psychology: Learning, Memory, and Cognition* 43, 11 (2017), 1699.
- [7] Cleeremans, Axel, and McClelland, James L. Learning the structure of event sequences. *Journal of Experimental Psychology: General* 120, 3 (1991), 235.

- [8] Clewett, David, DuBrow, Sarah, and Davachi, Lila. Transcending time in the brain: How event memories are constructed from experience. *Hippocampus* 29, 3 (2019), 162–183.
- [9] Cohen, Andrew L, Nosofsky, Robert M, and Zaki, Safa R. Category variability, exemplar similarity, and perceptual classification. *Memory & Cognition* 29, 8 (2001), 1165–1175.
- [10] Dayan, Peter. Improving generalization for temporal difference learning: The successor representation. *Neural computation* 5, 4 (1993), 613–624.
- [11] Dehaene, Stanislas, Meyniel, Florent, Wacongne, Catherine, Wang, Liping, and Pallier, Christophe. The neural representation of sequences: from transition probabilities to algebraic patterns and linguistic trees. *Neuron* 88, 1 (2015), 2–19.
- [12] Fiser, József, and Aslin, Richard N. Statistical learning of higher-order temporal structure from visual shape sequences. *Journal of Experimental Psychology: Learning, Memory, and Cognition* 28, 3 (2002), 458.
- [13] Gabay, Yafit, Dick, Frederic K, Zevin, Jason D, and Holt, Lori L. Incidental auditory category learning. *Journal of Experimental Psychology: Human Perception and Performance* 41, 4 (2015), 1124.
- [14] Gershman, Samuel J. The successor representation: its computational logic and neural substrates. *Journal of Neuroscience* 38, 33 (2018), 7193–7200.
- [15] Glenberg, Arthur M. Component-levels theory of the effects of spacing of repetitions on recall and recognition. *Memory & Cognition* 7, 2 (1979), 95–112.

- [16] Howard, Marc W, Fotedar, Mrigankka S, Datey, Aditya V, and Hasselmo, Michael E. The temporal context model in spatial navigation and relational learning: toward a common explanation of medial temporal lobe function across domains. *Psychological review* 112, 1 (2005), 75.
- [17] Kahn, Ari E, Karuza, Elisabeth A, Vettel, Jean M, and Bassett, Danielle S. Network constraints on learnability of probabilistic motor sequences. *Nature human behaviour* 2, 12 (2018), 936–947.
- [18] Karuza, Elisabeth A. The value of statistical learning to cognitive network science. *Topics in Cognitive Science* 14, 1 (2022), 78–92.
- [19] Karuza, Elisabeth A, Kahn, Ari E, and Bassett, Danielle S. Human sensitivity to community structure is robust to topological variation. *Complexity* 2019, 1 (2019), 8379321.
- [20] Karuza, Elisabeth A, Kahn, Ari E, Thompson-Schill, Sharon L, and Bassett, Danielle S. Process reveals structure: How a network is traversed mediates expectations about its architecture. *Scientific reports* 7, 1 (2017), 12733.
- [21] Knowlton, Barbara J, Ramus, Seth J, and Squire, Larry R. Intact artificial grammar learning in amnesia: Dissociation of classification learning and explicit memory for specific instances. *Psychological science* 3, 3 (1992), 172–179.
- [22] Kornell, Nate, and Bjork, Robert A. Learning concepts and categories: Is spacing the “enemy of induction”? *Psychological science* 19, 6 (2008), 585–592.
- [23] Kornell, Nate, Castel, Alan D, Eich, Teal S, and Bjork, Robert A. Spacing as the friend of both memory and induction in young and older adults. *Psychology and aging* 25, 2 (2010), 498.

- [24] Kruschke, John K. Alcove: An exemplar-based connectionist model of category learning. In *Connectionist Psychology*. Psychology Press, 2020, pp. 107–138.
- [25] Love, Bradley C. Comparing supervised and unsupervised category learning. *Psychonomic bulletin & review* 9, 4 (2002), 829–835.
- [26] Love, Bradley C, Medin, Douglas L, and Gureckis, Todd M. Sustain: a network model of category learning. *Psychological review* 111, 2 (2004), 309.
- [27] Lynn, Christopher W, and Bassett, Danielle S. How humans learn and represent networks. *Proceedings of the National Academy of Sciences* 117, 47 (2020), 29407–29415.
- [28] Lynn, Christopher W, Kahn, Ari E, Nyema, Nathaniel, and Bassett, Danielle S. Abstract representations of events arise from mental errors in learning and memory. *Nature communications* 11, 1 (2020), 2313.
- [29] Lynn, Christopher W, Papadopoulos, Lia, Kahn, Ari E, and Bassett, Danielle S. Human information processing in complex networks. *Nature Physics* 16, 9 (2020), 965–973.
- [30] Mack, Michael L, and Palmeri, Thomas J. The dynamics of categorization: Unraveling rapid categorization. *Journal of Experimental Psychology: General* 144, 3 (2015), 551.
- [31] Medin, Douglas L, and Bettger, Jeffrey G. Presentation order and recognition of categorically related examples. *Psychonomic bulletin & review* 1, 2 (1994), 250–254.
- [32] Medin, Douglas L, Wattenmaker, William D, and Hampson, Sarah E. Family resemblance, conceptual cohesiveness, and category construction. *Cognitive psychology* 19, 2 (1987), 242–279.

- [33] Momennejad, Ida, Russek, Evan M, Cheong, Jin H, Botvinick, Matthew M, Daw, Nathaniel Douglass, and Gershman, Samuel J. The successor representation in human reinforcement learning. *Nature human behaviour* 1, 9 (2017), 680–692.
- [34] Nissen, Mary Jo, and Bullemer, Peter. Attentional requirements of learning: Evidence from performance measures. *Cognitive psychology* 19, 1 (1987), 1–32.
- [35] Nosofsky, Robert M. Attention, similarity, and the identification–categorization relationship. *Journal of experimental psychology: General* 115, 1 (1986), 39.
- [36] Nosofsky, Robert M. The generalized context model: An exemplar model of classification. *Formal approaches in categorization* (2011), 18–39.
- [37] Nosofsky, Robert M, Gluck, Mark A, Palmeri, Thomas J, McKinley, Stephen C, and Glauthier, Paul. Comparing modes of rule-based classification learning: A replication and extension of shepard, hovland, and jenkins (1961). *Memory & cognition* 22, 3 (1994), 352–369.
- [38] Roark, Casey L, and Holt, Lori L. Task and distribution sampling affect auditory category learning. *Attention, Perception, & Psychophysics* 80 (2018), 1804–1822.
- [39] Roark, Casey L, Lehet, Matthew I, Dick, Frederic, and Holt, Lori L. The representational glue for incidental category learning is alignment with task-relevant behavior. *Journal of Experimental Psychology: Learning, Memory, and Cognition* 48, 6 (2022), 769.
- [40] Romberg, Alexa R, and Saffran, Jenny R. Statistical learning and language acquisition. *Wiley Interdisciplinary Reviews: Cognitive Science* 1, 6 (2010), 906–914.

- [41] Schapiro, Anna C, Rogers, Timothy T, Cordova, Natalia I, Turk-Browne, Nicholas B, and Botvinick, Matthew M. Neural representations of events arise from temporal community structure. *Nature neuroscience* 16, 4 (2013), 486–492.
- [42] Shepard, Roger N, Hovland, Carl I, and Jenkins, Herbert M. Learning and memorization of classifications. *Psychological monographs: General and applied* 75, 13 (1961), 1.
- [43] Stachenfeld, Kimberly L, Botvinick, Matthew M, and Gershman, Samuel J. The hippocampus as a predictive map. *Nature neuroscience* 20, 11 (2017), 1643–1653.
- [44] Turk-Browne, Nicholas B, Isola, Phillip J, Scholl, Brian J, and Treat, Teresa A. Multidimensional visual statistical learning. *Journal of Experimental Psychology: Learning, Memory, and Cognition* 34, 2 (2008), 399.
- [45] Unger, Layla, and Sloutsky, Vladimir M. Ready to learn: Incidental exposure fosters category learning. *Psychological Science* 33, 6 (2022), 999–1019.
- [46] Unger, Layla, Weichart, Emily R, Reardon, Noah, and Sloutsky, Vladimir. Without even trying: How incidental exposure shapes category learning. In *Proceedings of the Annual Meeting of the Cognitive Science Society* (2023), vol. 45.
- [47] Vlach, Haley A, Sandhofer, Catherine M, and Kornell, Nate. The spacing effect in children’s memory and category induction. *Cognition* 109, 1 (2008), 163–167.
- [48] Wahlheim, Christopher N, Dunlosky, John, and Jacoby, Larry L. Spacing enhances the learning of natural concepts: An investigation of mechanisms, metacognition, and aging. *Memory & cognition* 39 (2011), 750–763.
- [49] Whitehead, Peter S, Zamary, Amanda, and Marsh, Elizabeth J. Transfer of category learning to impoverished contexts. *Psychonomic Bulletin & Review* (2021), 1–10.

- [50] Zacks, Jeffrey M, and Swallow, Khena M. Event segmentation. *Current directions in psychological science* 16, 2 (2007), 80–84.